# The Decision No One Authored

The Answerability Gap in Generative AI

Luke F. Walton

Independent Researcher · lukefwalton.com

ORCID: 0009-0005-9263-1954

Preprint. Not yet peer-reviewed. · v2 · August 2026

Companion to *The Captured Oracle: Authorship and Agency in the Ethics of Answer-Engine Optimization*

https://doi.org/10.5281/zenodo.20676327

**Abstract.** Can a person retain full control of an AI-shaped decision and still fail to author it? Consider a hypothetical pension fund whose generative system guides its allocations through a decade of good returns. Officers review proposals, can override anything, and sign every trade: every outward mark of control, oversight, and authorship is present. Yet in this case, the conception of risk under which the money moves remains a vendor's default: no one at the fund ever decides what risk should mean for these beneficiaries. The decision appears authored, though the fund never made the judgment it turns on. The natural question is whether responsibility has shifted to the machine or remains secured by the humans around it. Debates over consciousness, patiency, and machine agency address the first possibility; meaningful human control and oversight address the second. Neither settles whether anyone performed that prior judgment. I call this the answerability gap. It is deepest where the system is best: reliability promotes the unowned frame into more consequential decisions and dissolves the checking that would expose it. The paper gives an account of the missing act and of the adoption that makes an inherited frame answerably one's own.

**Keywords:** answerability gap; moral patiency; attributability; moral responsibility; meaningful human control; automated decision-making

## Disclosures

**Competing interests.** The author is the founder of Surmado, Inc., which builds AI-orchestration systems for small businesses. The paper's diagnosis applies to that commercial work as much as to any other actor's. No funding was received for this work, and the analysis was neither commissioned nor reviewed by any commercial party. This work was conducted in the author's personal capacity; the views expressed are the author's own and do not represent those of any employer.

**Generative AI use.** Several frontier foundation models (from Anthropic, OpenAI, Google, and xAI) supported literature search, sectional drafting, argument pressure-testing, revision, and formatting. The author originated the thesis and its central distinctions, set the standards for inclusion, directed and revised all drafted material, and verified every claim, quotation, and citation against primary sources rather than against model agreement. The author is answerable for the final form.

## 1. The Decision and Its Subjects

Suppose a public pension fund adopts an AI system to guide its allocations. The system is built on a large language model: it reads filings, market data, and research, composes a fluent investment thesis for each position, and proposes how the portfolio should move. The proposals are tuned against a target the vendor ships as its default: long-run risk-adjusted return within the fund's liquidity constraints. The target names the optimization objective. It does not itself settle the conception of risk through which the system pursues it. An investment committee sets the mandate and reviews the record quarterly. Day to day, an analyst reads each thesis and signs each rebalancing before it executes, and can override the system at any point. After the fact the fund can say exactly which officer approved which trade. For ten years the system is excellent. Its proposals visibly and repeatedly outperform the committee's judgment, so the committee learns, as a rational body would, that its overrides subtract value. The committee's review thins. The committee widens the mandate. By most ordinary lights, human beings remain responsible for every allocation.

Now consider what has accumulated while nothing went wrong. The conception of risk inside the target is whatever a product built for no fund in particular came to carry. Measure and horizon, which correlations count, which tails are worth pricing, what may be hazarded and for whose sake, were not settings anyone installed. The vendor's engineers built a training system and fitted it to historical financial data. They are software engineers, not investment professionals, and would answer to a buyer for the fit, not for a view of risk. What the product measures was never settled as a judgment. What risk is to mean for these beneficiaries, under this mandate, is a question no product can settle. The fund never settled it. The officers who procured the system have moved on. The analysts who monitor it can say what the model weighs, and cannot say why that weighing is right for this fund and these beneficiaries. Each good year has widened what the frame governs and thinned the attention paid to it. The engineers composed a training system, the fund came to govern by what training left, and no one has ever made it the fund's own.

That is the first wrong, and it is done while every return is good: the fund has let consequential authority accumulate behind a frame no one at the fund can answer for. The beneficiaries whose retirements move under the system are decided about under a conception of risk that no party at the fund holds or would defend as its own; if one of them asks what the fund means by risk, the fund can produce a facsimile of a decision, yet cannot produce a decision, because none was made. What makes the lapse a wrong rather than an administrative gap is the direction of the authority. An institution that decides about other people's interests owes the people decided about more than good outcomes; it owes them terms that can be asked for, and an asking that finds a party who performed the judgment those terms express. Where no such judgment was performed, the answer owed is not withheld but missing, and the institution has kept the benefit of deciding while shedding the burden of having decided. That an institution owes the people it decides about a justification has a settled home in work on directed obligation and on the right to justification (Darwall 2006; Forst 2012), where what is owed is held against a party and answered to. What the fund's case adds is not a justification refused but a party who owes the account and has nothing to give. The judgment that account would cite is whether this conception of risk should govern this fund's decisions for these beneficiaries. It was never performed.

The second wrong waits on an edge. Some year a correlation regime shifts or a liquidity cascade arrives, the frame meets a case outside the historical training set, and the demand for reasons lands on the analyst who signed. She did not procure the system, configure its target, or choose its vendor, and she now stands answerable for the one thing in the chain that the fund never did. Holding her responsible will have the outward form of accountability and the inner structure of scapegoating. The inquiry will ask its natural question well, why the model erred that day. The prior question has no address: who decided that this conception of risk was the one under which billions were permitted to move. No one did.

The two wrongs share a single source. An act was owed and never occurred: the exercise of judgment over the decisions' evaluative frame, the settling of what the system is for and what its outputs must meet in order to be acted on.

Nothing in the case turns on the frame being wrong. Let the conception of risk be as good as such conceptions get, its estimates state of the art, the decade of returns earned and repeatable: both wrongs remain, because neither is a wrong of misjudgment.

Two literatures stand ready to say what has gone wrong here, and neither of them reaches it. The first is the debate over the machine's status. The second was built to preserve responsibility in machine-mediated action, and what it secures are standing relations: control, oversight, ownership. Each can hold completely in the fund's case. That is the gap this paper isolates, what I will call the answerability gap: a missing act, not a missing participant, and not a missing capacity either. No resolution of the machine's status closes the gap, and no strengthening of the human's standing conditions closes it, because acts are not entailed by standings.

## 2. The Status Question at Full Strength

The debate over the machine's status runs two questions together: patiency and answerability. By moral patiency I mean the property in virtue of which an entity can be wronged. For the welfare-based form of patiency at issue here, that property is a matter of whether there is something it is like to be it, in Nagel's (1974) phrase, such that what happens to it matters morally for its own sake. By answerability I mean the face of responsibility Shoemaker (2011) isolates under that name: an agent is answerable for an action insofar as it expresses her evaluative judgment, so that she is the apt target of a demand for justification and the fitting answer cites the judgment that was hers. The first is a question about the machine. The second is a question about whose evaluative judgment the action expresses.

The welfare case is stronger than its dismissals allow. Indicator-based assessment of current systems finds no strong evidence of consciousness, and their first-person reports, trained on human descriptions of experience and reversible by reprompting, are weak evidence either way (Butlin et al. 2023, 2026; Chalmers 2023; cf. Seth 2025). But absence of strong evidence leaves the question open. A growing literature argues that the uncertainty itself carries obligations: that near-future systems may be welfare subjects, that the probability cannot responsibly be rounded to zero, and that institutions should prepare for candidate moral patients rather than adjudicate them after the fact (Long et al. 2024; Birch 2024). Grant all of it. Grant, further, that the question may someday resolve in the affirmative, and that if it does, what we owe such systems will change in kind.

A second wing of the family concerns participation rather than experience. On the intentional-stance view, treating a system as a reasoner is a stance warranted by reasons-responsive competence, gradable and extensible in principle to anything competent enough (Dennett 1987). A capable model earns the stance, and refusing it a place in the space of reasons looks like an assumption rather than an argument. Grant the stance too, as far as competence carries it.

A third wing denies that the question is about intrinsic properties at all. Standing is conferred by how these systems come to figure in our practices (Coeckelbergh 2010; Gunkel 2018). Grant the reframing. And the family has a deflationary wing, which reaches the same place from the opposite verdict: present systems are artifacts, their apparent standing a design choice, and the ethics of their use therefore reduces to what humans owe humans (Bryson 2010).

The wings disagree about nearly everything, and a recurring assumption runs through the debate they compose: that the status verdict is what the ethics of use waits on. For the welfare wing, patiency arriving would rewrite that ethics from the ground up; for the participation wing, a participant machine enters the very practice responsibility lives in; for the relational wing, the conferral settles the standing; for the deflationary wing, the denial settles it the other way, and its formulation is the most revealing: current AI is probably not conscious, so humans remain in charge.

That sentence is hostage to its premise. It invites the rejoinder that the systems are becoming more agentic, more persistent, more plausibly candidates for some form of inner life, and that the ethics must therefore be revisited from the ground up as the premise weakens. The optimists' formulations are hostage in the mirror-image way: they treat the responsibility of the humans in the chain as provisional, pending the machine's promotion. The assumption is not universal, and its critics mark where this paper stands: Behdadi and Munthe (2020) argue that the debate over artificial moral agency should be redirected from metaphysical criteria to the normative question of how such systems should be included in practices that assume the agency and responsibility of their participants. What follows is a redirection of that kind, carried to the responsibility side and given its object: which act allocates the account for a machine-mediated decision, and whether anyone performed it. This paper concedes every empirical and metaphysical possibility the family contemplates, and denies the assumption that recurs across it. Who must answer for a machine-mediated decision is not settled by the machine's status, in either direction.

## 3. The Decoupling

Machine ethics has had the needed separation available for two decades. Floridi and Sanders (2004) argued that an artificial agent can be a genuine source of moral action without the mental states that would make it blameworthy, and treated agency and patiency as distinct roles an entity may occupy. The lesson is the structural separation, not their constructive program: being a source of moral action, being a fit bearer of responsibility, and being a possible recipient of moral treatment come apart. Status does not allocate the account for the frame. A patient machine need not be answerable, and an answerable machine would answer only for judgments it actually performed. A practice that confers standing on the machine has thereby made a decision, not discharged one.

The requirement is on the act, not on the performer. Suppose a system could recognize an evaluative choice as a choice, settle a conception of risk for a stated purpose, defend the settling under challenge and revise it, and stand in the practice in which such answers are owed. An institution might then delegate the setting of a frame to it explicitly, and nothing in this account forbids the delegation from succeeding: a party of any substrate that genuinely performed the setting would answer for the frame it set. What the machine's status and capabilities can change is who is available to perform the act and to answer for it. What they cannot do is settle, by themselves, whether the act was performed. Choosing to delegate the setting, and adopting what the delegate sets, remain judgments with the institution's name on them. In the fund's case the system before the analyst set nothing. It inherited a target, and no party of any substrate confronted it as the fund's.

Vallor and Vierkant (2024) converge on the human side of this from the responsibility-gap literature, arguing that current systems lack the reciprocal standing our responsibility practices require. The answerability at issue here is distinct from Tigard's (2021a) technological answerability, a system engineered to give answers on demand: saying why an output was produced leaves open who must answer for its being acted on.

## 4. The Missing Act

An evaluative frame is what a system is for and what its outputs must meet in order to be acted on; every deployed system operates under one. Setting a frame is making it answerably one's own. That is an act of judgment, and a frame can operate without that act. Someone confronts the value choice the frame encodes, recognizes it as a choice, and closes it, in someone's name. The act is occurrent, and its occurrence is a fact a record can show: a default examined, a target contested, a standard set where none was supplied, the closing locatable as someone's doing. What the record settles is allocative: it fixes in whose name the closing was done, so that the demand for justification has an address. It cannot certify that judgment stood behind the closing, since a check can be logged as performed and clicked past as reflexively as it was installed. The record says who must answer; whether there is an answer to give, only the act settles. Most frames are absorbed from somewhere, a profession, a tradition, a market, and an absorbed frame can still be made answerably one's own by adoption: a confrontation with the value choice that closes in the adopter's name and puts her in a practice that holds her to it. The setting cannot be discharged by procurement. A fund that buys a system carrying an operative conception of risk has transferred a product, not performed a judgment. The product is offered as if that transfer were the settling. A standing rule that the fund will accept whatever conception of risk a validated system embodies is that same transfer: it chooses a mechanism. The conception of risk for these beneficiaries is set only if someone in the chain confronts and closes it, or delegates the setting to a party that does. If no one makes that conception the fund's own, its fitness as the operative standard for these allocations has been confronted nowhere, and no quantity of ratification at the desk supplies the missing judgment.

The requirement is a reading of Shoemaker's condition. Answerability holds where the action expresses the agent's evaluative judgment, and an expression relation needs both of its terms: a judgment, and an agent who holds it. The allocations taken under the system's operative frame run on a determinate answer to what may be hazarded and for whose sake. That answer expresses no one's judgment until it is adopted. What the analyst holds is a view about the system's record, answerable in its own right and silent about risk. This is why the demand for justification, put to her about the frame, finds nothing of hers to cite in reply. Where a frame already expresses a judgment someone holds, no further ceremony is needed; where it expresses none, only the performing of one can make it do so, and the performing is an act.

The occurrent requirement is why a decision-maker can fail while passing every dispositional test.

Asked, the analyst would defend the frame, hear an objection, perhaps revise; but a standing willingness to justify what one merely received remains a disposition where an act was required. Reflective deference fails the same way: "I have considered it and I trust the system" makes the trust answerable while leaving the frame exactly as mute as she found it. A long tradition ties asserting to a standing readiness to defend what one has put forward (Brandom 1983). But that commitment runs to the product, the claim asserted, whereas the act at issue concerns who set the frame: one can be fully answerable for defending a recommendation while having set none of the choices that made it what it is.

Nothing bars the analyst from adopting the frame now. From the moment the value choice is confronted and closed in her name, later decisions are taken under someone's judgment. What the present act cannot do is reach backward. It does not supply the judgment under which the last decade's allocations were made. A standing willingness to endorse still leaves the act to be performed. Nor does an inherited conception restart the demand. If the analyst's view of risk is absorbed from her culture, the requirement seems to have no stop. The demand is instead whether anyone adopted this conception for these beneficiaries.

The frameworks built to preserve responsibility secure relations of a different kind, and their insufficiency has a single structure. The most developed is meaningful human control. It originates in the autonomous-weapons debate, where Sparrow (2007) argued that no candidate party, programmer, commander, or machine, could justly be held responsible for an atrocity. Santoni de Sio and van den Hoven (2018) gave the remedy philosophical foundations: a system is under meaningful human control when it tracks the relevant moral reasons of the relevant humans and traces to the appropriate moral understanding of at least one human agent in its design or use. The framework has since been operationalized for decision-support systems of exactly this kind (Cavalcante Siebert et al. 2023). The analyst satisfies it: she can override any proposal, the system tracks the fund's stated reasons, and its behavior traces to humans who understand it. Both conditions hold, and the failure has nonetheless occurred, because both conditions take the frame as given. Tracking is silent on whether the tracked reasons were the product of judgment or were inherited, unexamined; tracing reaches a human who understands the system, or who is merely in a position to understand it, without reaching one who has judged what its proposals should mean.

A defender will reply that tracking requires responsiveness to the relevant moral reasons, so a system that operates under an inherited conception of risk fails tracking after all. I grant the charitable reading. Tracking and tracing, however construed, are relational: they hold or fail between a system's behavior and the frame under which the system operates, and they say nothing about whether that frame was ever adopted. To fault the system for tracking a defective conception of risk already presupposes a better conception someone was answerable for setting. The notions come apart in both directions. A frame can be correct and yet unset: tracking and tracing can hold completely while no one has made the standard a prediction must meet their own judgment. Conversely, a frame can be set and yet wrong. Santoni de Sio and van den Hoven (2018, p. 11) note of a commander who knowingly deploys an autonomous weapon that cannot comply with the laws of armed conflict that "not only the tracing, but also the tracking condition is satisfied," even as the attack is unlawful and the commander culpable. Correctness is neither necessary nor sufficient for the act's having occurred.

Oversight mandates require capacities to interpret, to resist automation bias, and to override. A person can hold every capacity toward a frame she never set. Green (2022) shows the mandates often cannot even secure the capacities; the failure isolated here is the complementary one, where the capacities are real and the act is still missing. Decision-ownership requires that the decision-maker be positioned to endorse the values the system encodes (Zeiser 2024). Endorsement-capacity is a standing: it asks whether the agent would endorse the frame under which the system operates. Even the deflationary wing of the gap literature denies that machine-mediated action opens any novel gap at all (Tigard 2021b; Königs 2022; Demirtas 2025), and what it reaches is at most a locatable answerable party. Each of these is a standing, and the act is not entailed by it.

Matthias (2004) made canonical the attributive gap: as learning systems become more adaptive, no human may be connected to the outcome closely enough to bear responsibility at all. The answerability gap is that someone — often many people — is attributable for the outcome, and no one exercised the judgment over its evaluative frame that responsibility is supposed to track. The distinction is Shoemaker's (2011), here sorting gaps in a distributed act rather than kinds of one agent's responsibility. The fund assigns the outcome to the analyst. The appointment cannot make the judgment hers.

Kiener (2025) argues, against the gap tradition, that AI-mediated harms typically leave too many attributable parties rather than too few. Much of that abundance is engineered, and the count is orthogonal either way: a crowd of named parties can each have taken the frame ready-made from the next. Taking responsibility for a harm, as a normative power exercisable before or after it (Kiener 2022), can name a bearer after the fact. It cannot supply the judgment under which the allocations were made. Constantinescu and Kaptein (2025) map how responsibility for what is done with large language models should be distributed. Theirs is an answer to where responsibility should sit. It can succeed completely while the question pressed here stays open: what a distribution allocates are bearers and their standings, and no allocation of bearers entails that the evaluative judgment any of them is supposed to bear was ever performed.

An argument locates the requirement in the judging subject itself. Torrecilla-Pinero (2026) makes the limit anthropological: judgment is constitutively personal and therefore non-delegable, a constraint offered as a supplement to tracking and tracing, and the case that carries it, AI-steered coverage decisions in healthcare, is the fund's case in another sector. What is required here falls on the act: occurrent, pitched at the frame rather than the case before it, institutional as readily as personal, and open to a performer of any substrate.

The nearest ancestor of the requirement lies outside this literature. Fischer and Ravizza (1998) hold that an agent is responsible for what issues from a decisional mechanism only if she has taken responsibility for it. Their condition governs an agent's relation to her own mechanism, and it can be met in full by an analyst who has taken responsibility for how she decides and has never once confronted the conception of risk under which she signs. The gap falls nearest what Santoni de Sio and Mecacci (2021) call active responsibility, a position: whether an agent is placed so as to take responsibility going forward. The analyst is so placed. The answerability gap opens in the distance between being positioned to perform an act and the act's having been performed. A position is held; an act is done. Their four gaps concern who holds which position. This one concerns whether anything was done.

## 5. Where the Act Is Owed

The account distinguishes five junctures at which responsibility for a machine-mediated decision is discharged, or dropped. Two operate at the level of the frame, possibly far upstream of the person at the keyboard. A setting made there, once taken up, is inherited by every later output. Three recur with every output, bringing each one under the frame already in force: it is measured against the standard, admitted into action, and owned as it goes out. None of the five is a checklist entry; each is a place where the judgment required is evaluative rather than technical, and where the fluency of generative systems invites the judgment to be skipped while the outward form of a decision survives intact.

Grant the fund complete visibility into the model's weighing: every weight inspectable, the conception of risk readable as a specification. Visibility makes the frame available for judgment; the junctures are where that judgment is performed.

The first frame juncture is the ends. Every deployed system is for something, and what it is for is a value choice no output of the system can settle, since every output is produced under it. Whether an allocation system should maximize return within stated constraints, preserve intergenerational purchasing power, or weigh hazards the record has never priced is not a technical question. The choice hides inside an objective or a target variable, which makes a normative decision look like a modeling one; that translating an aim into a target is a discretionary, normatively loaded choice is documented in the problem-formulation literature (Passi and Barocas 2019). To settle the ends is to recognize the value choice and own it rather than accept the objective that arrives pre-installed. In these systems that objective often arrives as a training fit or in a system prompt.

The second is the standards: what counts as a good output, and by what criteria. A model can apply a standard, and it can generate a candidate rubric; the output cannot be its own sufficient warrant. A model may propose the criteria; a responsible party must examine them and adopt them rather than inherit them. A reward signal is that inheritance: it scores what counts as good under a criterion the deployer need not have set. Where standards are tacit, or read off the model's own confidence, this juncture has lapsed with every other safeguard in place.

The three per-output junctures follow. Verification reconnects an output to the reality it purports to be about: the claim checked against the source, the summary against the document, the recommendation against the case. What must be checked, and against what, is fixed with the standards — what verification adds is the performed check, that the checking occurred, on this case, as this party's act rather than an assumption. Acceptance is the moment a proposal becomes a deed: the draft becomes the sent email, the proposal the executed trade, the classification the entry in a permanent record. This is the hinge, where the human either exercises judgment or merely ratifies, and the deepest design failures collapse it by making acceptance the default or the path of least resistance. The limiting case is familiar from coding agents that offer auto-accept settings, running commands and writing files with no confirmation step, colloquially called "YOLO mode": the juncture removed by design. Release is the standing behind the artifact as it goes out, answerability for sending it and for what it does, including consequences unintended but foreseeable. It too is an act: the account is assumed when the artifact leaves in someone's name, before any harm and whether or not a question ever comes, so that a question, if it comes, finds a party already standing behind the deed.

Consider a physician using a system to draft the clinical note for an encounter. The end is one her practice has settled: an accurate record in the service of care, not throughput. The standards are external to the model and held by the profession: clinical accuracy, completeness, the norms of the record. Verification is built into the act, because she reads the draft against what happened in the room and corrects it. Acceptance is explicit: nothing enters the record until she signs. And release too is explicit: the final form goes out under her name, hers to answer for. The system does a great deal, perhaps producing language she would not have written unaided, yet nothing is missing, because its output remained a proposal, measured against standards it did not set, by the person who answers for the result.

Nor is the profession's role merely the analyst's inheritance one level up. The profession set its standards as standards, owned and revised as evaluative commitments, and the physician adopted them by an answerable act of her own, entering a practice that holds her to them. What remains open per encounter, the fit of the standard to the case before her, is exactly what her verification and acceptance answer for. This is why speed does not defeat her: the frame was set, occurrently, upstream, and she performs the per-output acts anew with each note, however fast she signs. Both the physician and the fund use AI heavily, and the difference is whether anyone ever made the frame their own.

Nor does the structure need billions to exist. Its most familiar instance is by now a stock figure: the hiring system that ranks applicants under an inherited definition of merit while a recruiter, free to override anything, works down the list. Unowned frames are already doing quiet work wherever fluent systems are adopted and trusted. What separates the mundane deployments from the fund is only the promotion: years of performance carrying such a frame into decisions an edge can make catastrophic. The gap is ordinary, and reliability is its ladder.

The everyday sense of authorship runs the other way. A model that drafts a paragraph has, in ordinary speech, authored it; nothing here denies that it made the thing. But making is execution, and execution transfers to a system without remainder; what does not transfer is answerability for the frame the artifact serves. The nearest neighbor in this region is Nyholm (2024), on whose account generative systems open gaps of meaning, authorship, and responsibility over their outputs, credit for an impressive output being harder to earn than blame for a harmful one. Each of those is a question about the product, and the question here is who set the frame the product serves. A system can supply every word and answerably author none of it, just as a person can answerably author a decision whose every word a system supplied. And when this paper says a frame was authored by no one, it means that no party answerably owned it, not that no party shaped it. A frame can be causally shaped, configured by a vendor, inherited from data, even gamed by an interested party who shapes what a channel makes salient, and still be owned by no one, because shaping what a channel makes salient is not yet authoring the evaluative claim it voices.[1]

Two objections meet the account at this point. The first says it is the romantic fiction of the solitary maker, when real institutions distribute responsibility across many hands. The objection mistakes distribution for dissolution. A film has hundreds of contributors and a structure of authorship nonetheless. The five junctures can be distributed across the people who set the ends, set the standards, verified the outputs, accepted them, and released them. The arrangement remains sound so long as each act was someone's answerable judgment and the chain can be reconstructed. Where groups are genuine agents, the bearer can be an institution (List 2021). What converts distribution into dissolution is the absence, at one or more junctures, of any hand at all. Nor does the account demand an individual behind every act. Where a body forms a judgment by a procedure, the judgment can be the body's without being any member's. Aggregation over reasons can commit an institution to a conclusion no member holds, and the institution can still be asked why, because the procedure was instituted, run, and closed in its name, and the closing leaves a record. That is the institutional form of the act, and it is what the fund never performed. The operative conception of

[1]Where an interested party does author the claim a channel voices, and conceals the hand behind it, the answerability is relocated rather than absent. That case, covert authorship on the answer channel, is outside this paper's scope.

risk passed through no procedure at all; it emerged from training and passed into use without being confronted as a judgment. Between a judgment that is the institution's though no member's, and a frame that is nobody's including the institution's, lies the whole distance this paper is measuring.

The second objection concerns demandingness: if the act must occur, much institutional life before AI fails the standard too, so either the gap is everywhere and the machine is incidental, or the standard quietly relaxes. The standard was always failable, and the institutional failure predates the technology by decades. The gap is not everywhere: adoption shows the standard ordinarily met, the physician satisfying all five junctures at working speed because her profession confronted the value choices answerably, upstream. And the machine matters: the standard is unchanged, and what generative fluency changes is the economics of failing it, the scale at which inherited frames are installed and the invisibility of their acceptance. The gap is old, and what is new is how cheaply it is produced and how little it shows.

The junctures are buildable. Generation can be separated from acceptance so that proposals stay legible as proposals until someone accepts them; standards can be held outside the model, where the model can be measured against them; the chain of who set what can be kept reconstructable across layers, and design research is taking up such commitments independently (Zhu et al. 2026). What architecture cannot do is manufacture the judgment. A target variable shown in bright letters can still be waved through, an inserted check clicked past as reflexively as it was added. Architecture can refuse to let the act's absence be invisible. It cannot perform the act for anyone.

## 6. The Question as Cover

If the act is ordinary and mostly met, why is it skipped here, and why does the skipping not show?

The first half of the answer is Vallor's (2024): contemporary AI is a mirror. Systems trained on the human archive do not stand outside our culture; they reflect our language, judgments, and institutional habits back to us, fluently enough for the reflection to be mistaken for an independent source of insight. The fluency operates with particular force at the point of judgment. A ranked list, a drafted paragraph, a recommendation: each arrives formatted as a conclusion, and the smoothness of the presentation is itself an argument for accepting it. The mirror does not merely tempt us to believe the machine understands; it tempts us to treat its outputs as if the evaluative work that would make them trustworthy were already done. Abdication therefore feels, from the inside, like reasonable reliance. Operators over-rely on automated aids they are positioned to monitor (Parasuraman and Riley 1997). The pattern persists in AI-assisted decision-making (Buçinca, Malaya, and Gajos 2021; Vasconcelos et al. 2023).

Capability sharpens the turn. Fluency tempts by making an output look finished; competence tempts by making deference correct. A reliably right system earns a trust the merely fluent one never could — the fit response to a track record — and it dissolves the felt need to check most where the system has performed best. The danger is no longer that the user mistakes a slick output for a sound one; it is that she stops looking, on good evidence, at the moment looking still mattered. The evidence is good and it is about the wrong quantity: how often an output is good under the frame in force, each entry produced under that frame and scored by it. What such a record cannot bear on is whether the frame is the one these decisions should be made under. Confidence well placed at every instance is compatible with a failure that accumulates across the whole run. Any system can produce the failure; a regression can carry an unowned frame as faithfully as a transformer. Generative fluency is what lets it scale and disappear: the more finished the output, the less visible the frame it inherits.

The second half of the answer is the one the public argument most easily misses about itself. The same fluency makes the question of whether anyone is home the natural question to ask about a system that talks like a someone. The status question is real, and it has been conceded everything. At the desk, it is also cover. Asking whether the machine really understands stands in for asking who set its terms, and the first question can be debated forever without the second being asked once. The analyst who wonders, idly, whether the system grasps what it is doing has already signed the rebalancing; the fund that convenes a working group on responsible AI can inherit the product's operative conception of risk unexamined for the length of the deliberation. The debate performs seriousness about the machine while the act lapses in silence, and the better the system, the more natural the wrong question becomes. Every increment of fluency is an increment in the apparent urgency of asking what the machine is, and none of it bears on who set the frame. Arendt (1963) named the local failure thoughtlessness: action continued without the reflective interruption in which one asks what one is actually doing. What transfers is the concept, not the case. A person can be present, capable, and procedurally compliant while the evaluative judgment the action turns on was never performed. The interruption the act requires asks a different question from the one the machine's fluency invites. Not: is anything home? But: what is being done here, and who answers for it?

## 7. The Ratchet

When the account for a decision taken under an unowned frame comes due, it does not come due evenly. Elish's (2019) moral crumple zone names the operator faulted despite too little control. Here the operator has control in abundance, and what is missing is judgment, because the competence that makes intervention unnecessary also makes abdication invisible. And the zone is assigned by design. A crumple zone is engineered into the cheapest, most replaceable region of the structure.

The party on whom the account settles is selected by the same institutional logic that left the frame unowned: the front-line approver, nearest the moment a proposal becomes a deed, most junior, most replaceable, least able to refuse the position or push it back upward. The analyst did not procure the system, configure its target, or choose the vendor; she is the last hand to touch the decision and the first the institution can afford to lose.

The problem of many hands (Thompson 1980) is usually told as diffusion, and the diffusion is asymmetric. It strips answerability from the parties with the standing to shed it and deposits the remainder on the party with the least — agency laundering (Rubel, Castro, and Pham 2019), then the laundered account coming to rest. “Human in the loop” names a safeguard, and it is also where the account comes to rest, an occupant where an act was owed. The account does not vanish into the crowd of hands. The answerability gap is not dangerous because blame goes nowhere. It is dangerous because blame goes somewhere predictable: downhill.

The rotation has a second effect, on the frame rather than the person, and it is the more lasting. When the analyst leaves and another fills the chair, the unowned frame is handed on intact, and the next occupant inherits it exactly as she did, as a given, not a choice. Across enough cycles the frame is continuously unowned. It drifts, since the inherited system can change while no one reopens the frame. Each occupant ratifies what the last left in place, and at no point does any party confront the drift as a decision. Vaughan’s (1996) study of the Challenger launch is the pattern at institutional scale, before any generative system existed: the solid rocket boosters’ O-rings eroded on flight after flight, beyond anything their designers had predicted, and because no launch failed, the erosion was carried through the formal reviews and recorded there as an acceptable risk, a normalization of deviance. Repeated success can normalize a standard without anyone reopening the evaluative question it answers. A competent system industrializes the normalization of the gap, one user and one uneventful Tuesday at a time.

And the record of intact returns does more than quiet the checking — it is a credential. A frame that performs is promoted on the strength of its record and handed decisions of greater and greater consequence, so that exposure grows exactly as scrutiny shrinks. The unowned frame arrives at the decisions least able to bear it by having been right about the decisions that could. The risk this stores differs in kind from the risk of failure: it is constituted by the missing act, and accuracy deepens it, since accuracy is what earns the promotion and retires the check.

The fund now runs on a conception of risk no one there made its own. What no one can give is why this conception should govern these beneficiaries. Officers remain positioned to close it, and every year of good returns raises the price of that closing and thins the baseline it would stand on. The gap is a ratchet.

The failure is institutional before it is personal. The ends and standards were already in force upstream of the analyst, in what procurement took up and what training left, and that is where the missing act was owed. Blaming the approver who inherited the frame is itself a form of the laundering described above. The erosion is old in kind: Nissenbaum (1996) traced how computerized systems obscure accountability, through many hands, through the diffusion of fault in software, and through the temptation to treat the computer itself as the answerable party. What generative fluency adds is scale and silence. Run the ratchet forward across an economy of deployments and the aggregate comes into view: an administrative order in which an ever-growing share of consequential decisions passes through frames that are nobody's, revised by no one because no one holds a baseline a revision could state and stand behind, contested by no one because their subjects never meet them at all, until someone answers for everything and authored none of it.

## 8. Conclusion

Whether these systems are, or might become, minds we could wrong is a real question, and I have not settled it; I have argued that it is the wrong question to put at the center of the ethics of using them, because its answer, whatever it proves to be, settles nothing by itself about who performed the judgment a given decision turned on. Patiency is uncertain and may change; the question of who answers is immediate, and it is asked act by act. We can usually name the human who signed off. What we cannot assume is that she authored what she signed.

The deepest temptation the mirror presents is not that we will mistake the machine for a person, but that we will let it do our judging and call the result a decision. Fluency opens that door and competence walks us through it: reliability earns a trust that makes not-checking reasonable, and the reasonable abdication is the hardest kind to interrupt, because nothing about it feels like surrender. The discipline this requires is the reflective interruption described earlier: the moment in which a person asks what is actually being done, and whether she can answer for it. To build and use these systems well is to keep that interruption alive, and to keep the sites of it reconstructable: ends and standards made someone's, acceptance that remains an act, so that someone remained answerable for the judgment the decision turned on. Much of what a system does — retrieval, formatting, computation, first-pass synthesis — can be handed over entirely, provided those sites remain someone's. Neglecting the answerability gap extends the separation of liability from authorship, until someone answers for everything and authored none of it. The likelihood of such a future increases as the machines improve. Every increment of capability is an increment in the rational case for deference, and so in the ease of the abdication: the better the system, the less anything feels wrong as the judgment quietly stops being anyone's.

## References

Arendt, H. (1963). *Eichmann in Jerusalem: A report on the banality of evil.* Viking Press.

Behdadi, D., & Munthe, C. (2020). A normative approach to artificial moral agency. *Minds and Machines, 30*(2), 195–218. https://doi.org/10.1007/s11023-020-09525-8

Birch, J. (2024). *The edge of sentience: Risk and precaution in humans, other animals, and AI.* Oxford University Press.

Brandom, R. (1983). Asserting. *Noûs, 17*(4), 637–650. https://doi.org/10.2307/2215086

Bryson, J. J. (2010). Robots should be slaves. In Y. Wilks (Ed.), *Close engagements with artificial companions: Key social, psychological, ethical and design issues* (pp. 63–74). John Benjamins.

Buçinca, Z., Malaya, M. B., & Gajos, K. Z. (2021). To trust or to think: Cognitive forcing functions can reduce overreliance on AI in AI-assisted decision-making. *Proceedings of the ACM on Human-Computer Interaction, 5*(CSCW1), Article 188. https://doi.org/10.1145/3449287

Butlin, P., Long, R., Bayne, T., Bengio, Y., Birch, J., Chalmers, D., Constant, A., Deane, G., Elmoznino, E., Fleming, S. M., Ji, X., Kanai, R., Klein, C., Lindsay, G., Michel, M., Mudrik, L., Peters, M. A. K., Schwitzgebel, E., Simon, J., & VanRullen, R. (2026). Identifying indicators of consciousness in AI systems. *Trends in Cognitive Sciences, 30*(6), 488–501. https://doi.org/10.1016/j.tics.2025.10.011

Butlin, P., Long, R., Elmoznino, E., Bengio, Y., Birch, J., Constant, A., Deane, G., Fleming, S. M., Frith, C., Ji, X., Kanai, R., Klein, C., Lindsay, G., Michel, M., Mudrik, L., Peters, M. A. K., Schwitzgebel, E., Simon, J., & VanRullen, R. (2023). Consciousness in artificial intelligence: Insights from the science of consciousness. arXiv. https://doi.org/10.48550/arXiv.2308.08708

Cavalcante Siebert, L., Lupetti, M. L., Aizenberg, E., Beckers, N., Zgonnikov, A., Veluwenkamp, H., Abbink, D., Giaccardi, E., Houben, G.-J., Jonker, C. M., van den Hoven, J., Forster, D., & Lagendijk, R. L. (2023). Meaningful human control: Actionable properties for AI system development. *AI and Ethics, 3*(1), 241–255. https://doi.org/10.1007/s43681-022-00167-3

Chalmers, D. J. (2023, August 9). Could a large language model be conscious? *Boston Review.* https://www.bostonreview.net/articles/could-a-large-language-model-be-conscious/

Coeckelbergh, M. (2010). Robot rights? Towards a social-relational justification of moral consideration. *Ethics and Information Technology, 12*(3), 209–221. https://doi.org/10.1007/s10676-010-9235-5

Constantinescu, M., & Kaptein, M. (2025). Responsibility gaps, LLMs & organisations: Many agents, many levels, and many interactions. *Science and Engineering Ethics, 31*, Article 36.

https://doi.org/10.1007/s11948-025-00560-1

Darwall, S. (2006). *The second-person standpoint: Morality, respect, and accountability.* Harvard University Press.

Demirtas, H. (2025). AI responsibility gap: Not new, inevitable, unproblematic. *Ethics and Information Technology, 27*, 7. https://doi.org/10.1007/s10676-024-09814-1

Dennett, D. C. (1987). *The intentional stance.* MIT Press.

Elish, M. C. (2019). Moral crumple zones: Cautionary tales in human-robot interaction. *Engaging Science, Technology, and Society, 5*, 40–60. https://doi.org/10.17351/ests2019.260

Fischer, J. M., & Ravizza, M. (1998). *Responsibility and control: A theory of moral responsibility.* Cambridge University Press.

Floridi, L., & Sanders, J. W. (2004). On the morality of artificial agents. *Minds and Machines, 14*(3), 349–379. https://doi.org/10.1023/B:MIND.0000035461.63578.9d

Forst, R. (2012). *The right to justification: Elements of a constructivist theory of justice* (J. Flynn, Trans.). Columbia University Press. (Original work published 2007)

Green, B. (2022). The flaws of policies requiring human oversight of government algorithms. *Computer Law & Security Review, 45*, Article 105681. https://doi.org/10.1016/j.clsr.2022.105681

Gunkel, D. J. (2018). *Robot rights.* MIT Press.

Kiener, M. (2022). Can we bridge AI's responsibility gap at will? *Ethical Theory and Moral Practice, 25*(4), 575–593. https://doi.org/10.1007/s10677-022-10313-9

Kiener, M. (2025). AI and responsibility: No gap, but abundance. *Journal of Applied Philosophy, 42*(1), 357–374. https://doi.org/10.1111/japp.12765

Königs, P. (2022). Artificial intelligence and responsibility gaps: What is the problem? *Ethics and Information Technology, 24*(3), Article 36. https://doi.org/10.1007/s10676-022-09643-0

List, C. (2021). Group agency and artificial intelligence. *Philosophy & Technology, 34*(4), 1213–1242. https://doi.org/10.1007/s13347-021-00454-7

Long, R., Sebo, J., Butlin, P., Finlinson, K., Fish, K., Harding, J., Pfau, J., Sims, T., Birch, J., & Chalmers, D. (2024). Taking AI welfare seriously. arXiv. https://doi.org/10.48550/arXiv.2411.00986

Matthias, A. (2004). The responsibility gap: Ascribing responsibility for the actions of learning automata. *Ethics and Information Technology, 6*(3), 175–183. https://doi.org/10.1007/s10676-004-3422-1

Nagel, T. (1974). What is it like to be a bat? *The Philosophical Review, 83*(4), 435–450. https://doi.org/10.2307/2183914

Nissenbaum, H. (1996). Accountability in a computerized society. *Science and Engineering Ethics, 2*(1), 25–42. https://doi.org/10.1007/BF02639315

Nyholm, S. (2024). Generative AI's gappiness: Meaningfulness, authorship, and the credit-blame asymmetry. In A. Strasser (Ed.), *Anna's AI anthology: How to live with smart machines?* (pp. 167–194). Xenomoi Verlag.

Parasuraman, R., & Riley, V. (1997). Humans and automation: Use, misuse, disuse, abuse. *Human Factors, 39*(2), 230–253. https://doi.org/10.1518/001872097778543886

Passi, S., & Barocas, S. (2019). Problem formulation and fairness. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* '19)* (pp. 39–48). Association for Computing Machinery. https://doi.org/10.1145/3287560.3287567

Rubel, A., Castro, C., & Pham, A. (2019). Agency laundering and information technologies. *Ethical Theory and Moral Practice, 22*(4), 1017–1041. https://doi.org/10.1007/s10677-019-10030-w

Santoni de Sio, F., & Mecacci, G. (2021). Four responsibility gaps with artificial intelligence: Why they matter and how to address them. *Philosophy & Technology, 34*(4), 1057–1084. https://doi.org/10.1007/s13347-021-00450-x

Santoni de Sio, F., & van den Hoven, J. (2018). Meaningful human control over autonomous systems: A philosophical account. *Frontiers in Robotics and AI, 5*, Article 15. https://doi.org/10.3389/frobt.2018.00015

Seth, A. K. (2025). Conscious artificial intelligence and biological naturalism. *Behavioral and Brain Sciences.* Advance online publication. https://doi.org/10.1017/S0140525X25000032

Shoemaker, D. (2011). Attributability, answerability, and accountability: Toward a wider theory of moral responsibility. *Ethics, 121*(3), 602–632. https://doi.org/10.1086/659003

Sparrow, R. (2007). Killer robots. *Journal of Applied Philosophy, 24*(1), 62–77. https://doi.org/10.1111/j.1468-5930.2007.00346.x

Thompson, D. F. (1980). The moral responsibility of public officials: The problem of many hands. *American Political Science Review, 74*(4), 905–916. https://doi.org/10.2307/1954312

Tigard, D. W. (2021a). Technological answerability and the severance problem: Staying connected by demanding answers. *Science and Engineering Ethics, 27*, Article 59. https://doi.org/10.1007/s11948-021-00334-5

Tigard, D. W. (2021b). There is no techno-responsibility gap. *Philosophy & Technology, 34*(3), 589–607. https://doi.org/10.1007/s13347-020-00414-7

Torrecilla-Pinero, J. A. (2026). Judgment cannot be delegated: A subject-preserving framework for AI governance. *Philosophy & Technology, 39*(3), Article 135. https://doi.org/10.1007/s13347-026-01140-2

Vallor, S. (2024). *The AI mirror: How to reclaim our humanity in an age of machine thinking.* Oxford University Press.

Vallor, S., & Vierkant, T. (2024). Find the gap: AI, responsible agency and vulnerability. *Minds and Machines, 34*(3), Article 20. https://doi.org/10.1007/s11023-024-09674-0

Vasconcelos, H., Jörke, M., Grunde-McLaughlin, M., Gerstenberg, T., Bernstein, M. S., & Krishna, R. (2023). Explanations can reduce overreliance on AI systems during decision-making. *Proceedings of the ACM on Human-Computer Interaction, 7*(CSCW1), Article 129. https://doi.org/10.1145/3579605

Vaughan, D. (1996). *The Challenger launch decision: Risky technology, culture, and deviance at NASA.* University of Chicago Press.

Zeiser, J. (2024). Owning decisions: AI decision-support and the attributability-gap. *Science and Engineering Ethics, 30*, Article 27. https://doi.org/10.1007/s11948-024-00485-1

Zhu, L., Lu, Q., Ding, M., Lee, S. U., & Wang, C. (2026). Designing meaningful human oversight in AI. *AI and Ethics, 6*, Article 286. https://doi.org/10.1007/s43681-026-01147-7